MoneyMinds — AI Analysis Simply Explained

For: Platform Owners & Clients **Updated:** November 2026 **Prepared by:** MoneyMinds Development Team

---

What Does "AI" Actually Do on This Platform?

Imagine you hired a **very fast researcher** who:

- Reads someone's name and profile
- Searches Google for them in seconds
- Reads all the results
- Writes a full report for you

That is exactly what our AI does — but it does it in under **30 seconds**, for any user, any time, automatically.

What Is a "Token"?

When the AI reads or writes something, it counts in small pieces called **tokens**.

> ***Simple way to think about it:*** *A token is like a* ***small word or piece of a word****. "Hello" = 1 token. A full paragraph = about 100 tokens.*

The platform pays for how much the AI **reads** (input) and how much it **writes back** (output).

What you send to AI	Approx. size
One sentence	20–30 tokens
One paragraph	100–150 tokens
One full page	400–500 tokens
Full search results	2,000–3,000 tokens

Two Tools Working Together

Every analysis uses **two tools** working as a team:

Tool 1 — Tavily (The Researcher)

Tavily searches the internet — like Google, but designed for AI.

- It finds real web pages, news, LinkedIn profiles, company info
- It charges per **search** (not per token)
- **Cost: about $0.005 per search** (half a cent)

Tool 2 — Groq AI (The Writer)

Groq reads all the search results and writes the final report.

- It is **10–20x faster** than regular AI services
- It charges per token (reading + writing)
- **Currently on free plan** — costs us $0

What Happens During Each Analysis?

Background Check on a New User

This is the most important analysis — it runs when admin clicks "Run Verification."

```
User applies → AI searches internet → AI reads results → AI writes report
   (5 sec)         (10 sec, 3–6         (reads ~3,000      (writes full
                    searches)             tokens)            report in
                                                             ~2,000 words)
```

What the report includes:

- Is this person real? (LinkedIn found or not)

- Do they work where they say they do?
- Any news about fraud or legal cases?
- Overall recommendation: Accept / Review / Reject

Cost per check: ~$0.02 to $0.05 (less than one rupee)

Investment Research Report

When an investor asks for a market analysis:

- AI searches for latest market news (1 search)
- Reads the results
- Writes a full report with opportunities, risks, and signals

Cost per report: ~$0.01 (almost nothing)

How Much Does It Cost Overall?

Situation	Monthly Cost
50 users, light usage	~$1–3
100 users, regular usage	~$3–8
500 users, heavy usage	~$15–30

> ***Right now, at current scale, the AI costs roughly the same as a cup of chai per month.***

Which AI Are We Using?

We use **three AI models** depending on the task. Think of them like different staff members — each hired for a specific job.

Staff Member 1 — The Main Analyst

Model: GPT-20B (served by Groq)

This is the AI that does most of the work — reading search results, writing reports, analyzing signals.

Feature	Detail
Speed	Very fast (results in 2–5 seconds)
Quality	Good — suitable for business reports
Current cost	Free (on Groq's free plan)
Best for	Background checks, signal analysis, reports

Staff Member 2 — The Document Reader

Model: Qwen Vision (served by Groq)

This AI is specially trained to **read images and documents** — like uploaded certificates, degrees, or charts.

Feature	Detail
Special skill	Can read text inside images and photos
Speed	Very fast
Current cost	Free
Best for	Certificate verification, document uploads

Staff Member 3 — The Backup (Claude by Anthropic)

Model: Claude Haiku (by Anthropic)

This is our **backup AI** — if Groq is ever unavailable, the platform automatically switches to Claude. Claude is made by Anthropic, one of the most trusted AI safety companies in the world.

Feature	Detail
Reliability	Extremely high — enterprise grade
Quality	Excellent
Cost	~$2–3 per month at current usage
Best for	Backup, high-accuracy tasks

> ***Fun fact:*** *Claude is also the AI that was used to help build MoneyMinds itself.*

Can We Upgrade the AI Later?

Yes — and it is very easy. **No code changes needed.** Just update one setting.

Here is the upgrade path as the platform grows:

Stage	AI Used	Quality	Monthly Cost
Now (Launch)	GPT-20B on Groq	Good	$0
Growth (100+ users)	Llama 70B on Groq	Very Good	$0 (still free)
Scale (500+ users)	Claude Sonnet	Excellent	~$15–20
Enterprise (1000+ users)	Claude Opus	Best available	~$75–100

The system was built to **start free and scale affordably** — the client never needs to pay for AI until the platform has hundreds of active users.

Why Groq? Why Not ChatGPT or Others?

Provider	Speed	Cost	Quality
Groq (what we use)	⚡ Fastest	💚 Free/very cheap	✅ Good
OpenAI (ChatGPT)	Medium	💰 Expensive	✅ Excellent
Google Gemini	Fast	🧡 Medium	✅ Good
Anthropic Claude	Fast	🧡 Medium	✅ Excellent

> ***We chose Groq because speed matters for user experience****, and their free tier is generous enough to run the platform at no AI cost during the launch phase. When quality becomes the priority over cost, we switch to Claude — one line change.*

Summary in Plain English

1. **The AI is like a fast researcher** — it searches, reads, and writes reports automatically.
2. **We use three AI tools** — a main analyst, a document reader, and a backup. Each does its job without the user seeing any of it.
3. **Current cost is essentially zero** — Groq's free plan handles our current scale comfortably.
4. **Tavily (web search) costs a small amount** — about $2–5 per month for 100 users.
5. **We can upgrade any time** — better AI is available at a few dollars per month. The upgrade requires changing one setting, not rewriting the app.
6. **The platform gets smarter as profiles are more complete** — the more details a user fills in (LinkedIn, university, company), the more accurate and useful the AI reports become.

Document prepared by MoneyMinds Development Team — November 2026 For client and stakeholder communication. No technical background required.